

JAMES ZENG

+1 (872) 258-8789 | San Francisco, CA (Open to relocate) | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) | james1.zeng@gmail.com

WORK EXPERIENCE

Future Intelligence Labs - AI Engineer San Francisco/CA

12/2025 - Present

- Built an autonomous market-data platform for a tire retailer (14% U.S. unit sales), replacing monthly buys prep with daily delivery.
- Designed agents to detect anomalies and choose reruns, repairs, or isolation across 12K+ nightly jobs, reducing routine intervention.
- Orchestrated agent fleets across AWS Lambda, cutting overnight recovery from 9 hours to 20 minutes and compute cost 83%.
- Partitioned collection workloads by 21 states, sustaining 2K jobs/hour while meeting a 3 hours overnight freshness window.
- Built *AgentTrace*'s([link](#)) pipeline, load-tested to ~100K events/sec, and React dashboard to correlate and visualize agent workflows.
- Enabled security teams to inspect workflows across languages, frameworks, and shadow agents without SDKs or app-team redeloys.
- Fixed collector reattachment after pod restarts, preventing a silent failure mode that lost ~87% of traces before storage.

Tech: TypeScript, Go, React, Next.js, Python, FastAPI, AWS EC2/Lambda/RDS/SES, GCP Pub/Sub, Kubernetes, agent orchestration

Berkeley SkyDeck - Full-Stack Intern Emeryville/CA

04 - 07/2025

- Built PromptArena([link](#)) for startup program owners, replacing subjective prompt evals with evidence-gated promotion to production.
- Implemented pairwise automated graders and bootstrap confidence interval to analyze LLM quality and benchmark inconsistencies.
- Added deterministic canaries, drift alerts, and regression tests; stress-tested 486 concurrent workflows with zero observed failures.
- Built and maintained PyTorch/TypeScript evaluation infrastructure, while cutting 1M-row result rendering from 15s to ~1s.

Tech: Python, PyTorch, TypeScript, React, LLM evaluation, benchmark development, automated testing, WebSockets, Docker, CI/CD

EndoVas Medical - Software Engineer Intern remote & Shanghai/CN

09/2023 - 07/2024

- Ran 11 onsite validation demos, reproduced failures with users, and translated communication into product and operational changes.
- Automated device setup and validation with port detection and prefilled configurations, replacing error-prone manual parameter entry.
- Shipped a FastAPI backend for remote medical-device monitoring with OAuth/JWT auth and automated software updates.
- Improved AWS Lambda analytics pipelines, cutting device troubleshooting time from 3 hours to 30 minutes.

Tech: Python, FastAPI, PostgreSQL, OAuth/JWT, AWS Lambda

PROJECT EXPERIENCE

Branch Master - Open-Source Project San Francisco/CA

05 - 07/2026

- Built a live branch map across multiple AI chat platforms, including ChatGPT, Claude, and Grok, to visualize and revisit hidden paths.
- Created site-specific adapters for streaming DOMs and automated native branch controls without sending or modifying messages.
- Designed local-first summaries that keep conversation data on-device with zero network access.
- Built 55 Chromium E2E checks covering streaming, edits, branch navigation, and zero-network privacy.

J.P. Morgan Chase Hackathon - Lead Developer (1st Place) Columbus/OH

03 - 04/2024

- Led solution architecture and core implementation as team lead and primary developer, placing 1st among 312 teams.
- Built a Scikit-Learn recommender that improved MRR 35% over baseline and a live dashboard that cut analytics latency ~20%.

EDUCATION

University of California, Berkeley - Master's Analytics, GPA: 3.88

2025

Ohio State University - Bachelor's Statistics; Minor in Computer Science, GPA 3.81

2024

SKILLS

AI & Agents: Agent orchestration, tool calling/MCP, LangGraph/LangChain, RAG, evaluation, observability & tracing

Engineering: Python, Go, TypeScript, React/Next.js, HTML, FastAPI, PostgreSQL, MongoDB, WebSockets

Infrastructure: AWS, GCP Pub/Sub, Docker, Kubernetes, CI/CD, distributed systems